

STAT 231: Problem Set 6B

Sean Wei

due by 5 PM on Friday, October 9

This homework assignment is designed to help you further ingest, practice, and expand upon the material covered in class over the past week(s). You are encouraged to work with other students, but all code and text must be written by you, and you must indicate below who you discussed the assignment with (if anyone).

Steps to proceed:

1. In RStudio, go to File > Open Project, navigate to the folder with the course-content repo, select the course-content project (course-content.Rproj), and click "Open"
2. Pull the course-content repo (e.g. using the blue-ish down arrow in the Git tab in upper right window)
3. Copy ps6B.Rmd from the course repo to your repo (see page 6 of the GitHub Classroom Guide for Stat231 if needed)
4. Close the course-content repo project in RStudio
5. Open YOUR repo project in RStudio
6. In the ps6B.Rmd file in YOUR repo, replace "YOUR NAME HERE" with your name
7. Add in your responses, committing and pushing to YOUR repo in appropriate places along the way
8. Run "Knit PDF"
9. Upload the pdf to Gradescope. Don't forget to select which of your pages are associated with each problem. *You will not get credit for work on unassigned pages (e.g., if you only selected the first page but your solution spans two pages, you would lose points for any part on the second page that the grader can't see).*

If you discussed this assignment with any of your peers, please list who here:

ANSWER:

Trump Tweets

David Robinson, Chief Data Scientist at DataCamp, wrote a blog post “Text analysis of Trump’s tweets confirms he writes only the (angrier) Android half”.

He provides a dataset with over 1,500 tweets from the account `realDonaldTrump` between 12/14/2015 and 8/8/2016. We’ll use this dataset to explore the tweeting behavior of `realDonaldTrump` during this time period.

First, read in the file. Note that there is a `TwitterR` package which provides an interface to the Twitter web API. We’ll use this R dataset David created using that package so that you don’t have to set up Twitter authentication.

```
load(url("http://varianceexplained.org/files/trump_tweets_df.rda"))
```

A little wrangling to warm-up

1a. There are a number of variables in the dataset we won’t need.

- First, confirm that all the observations in the dataset are from the screen-name `realDonaldTrump`.
- Then, create a new dataset called `tweets` that only includes the following variables:
- `text`
- `created`
- `statusSource`

```
# Confirm that realDonaldTrump is the only screen name in the dataset  
unique(trump_tweets_df$screenName)
```

```
## [1] "realDonaldTrump"
```

```
# New dataset with only text, created, and statusSource  
tweets <- trump_tweets_df %>%  
  select(text, created, statusSource)  
glimpse(tweets)
```

```
## Rows: 1,512  
## Columns: 3  
## $ text      <chr> "My economic policy speech will be carried live at 12:...  
## $ created   <dtm> 2016-08-08 15:20:44, 2016-08-08 13:28:20, 2016-08-08 ...  
## $ statusSource <chr> "<a href=\"http://twitter.com/download/android\" rel=\\"...
```

1b. Using the `statusSource` variable, compute the number of tweets from each source. How many different sources are there? How often are each used?

ANSWER: There are 5 unique tweet sources: Instagram, Twitter for Android, Twitter for iPad, Twitter for iPhone, and Twitter Web Client. There were 762 tweets from Twitter for Android, 628 tweets from Twitter for iPhone, 120 tweets from Twitter Web Client, and 1 tweet from both Instagram and Twitter for iPad.

```
# Isolate the text in the anchor tags
```

```
tweet_sources <- tweets %>%  
  mutate(  
    statusSource = gsub("</a>", "", statusSource),  
    statusSource = gsub(".*>", "", statusSource)  
  )  
glimpse(tweet_sources)
```

```
## Rows: 1,512
```

```
## Columns: 3
```

```
## $ text      <chr> "My economic policy speech will be carried live at 12:...
```

```
## $ created   <dtm> 2016-08-08 15:20:44, 2016-08-08 13:28:20, 2016-08-08 ...
```

```
## $ statusSource <chr> "Twitter for Android", "Twitter for iPhone", "Twitter ...
```

```
# Put each observation in statusSource into one big string separated by commas
```

```
source_count <- paste(tweet_sources$statusSource, collapse = ",")
```

```
# Turn large string into a character array
```

```
source_count <- strsplit(source_count, ",")
```

```
# Create frequency table for the character array
```

```
table(source_count)
```

```
## source_count
```

```
##      Instagram Twitter for Android  Twitter for iPad  Twitter for iPhone
```

```
##              1              762              1              628
```

```
## Twitter Web Client
```

```
##              120
```

1c. We're going to compare the language used between the Android and iPhone sources, so only want to keep tweets coming from those sources. Explain what the `extract` function (from the `tidyverse` package) is doing below. (Note that "regex" stands for "regular expression".)

ANSWER: The `extract` function is given a regular expression to find groups and then makes a new column based on them. Below, the `extract` function looks for instances that begin with "Twitter for" and then makes a new column with values of what comes after it and before the `<`. If the `statusSource` doesn't match this regex, the value of the new column is N/A. For example, if the `statusSource` was "...Twitter for Android<," the value in the new source column would be "Android," but if the `statusSource` was "...Twitter Web Client<," the value in the new source column would be "N/A."

```
tweets2 <- tweets %>%
  extract(col = statusSource, into = "source"
    , regex = "Twitter for (.*)<"
    , remove = FALSE) %>%
  filter(source %in% c("Android", "iPhone"))
glimpse(tweets2)
```

```
## Rows: 1,390
## Columns: 4
## $ text      <chr> "My economic policy speech will be carried live at 12:...
## $ created   <dtm> 2016-08-08 15:20:44, 2016-08-08 13:28:20, 2016-08-08 ...
## $ statusSource <chr> "<a href=\"http://twitter.com/download/android\" rel=\...
## $ source     <chr> "Android", "iPhone", "iPhone", "Android", "Android", "..."
```

How does the language of the tweets differ by source?

2a. Create a word cloud for the top 50 words used in tweets sent from the Android. Create a second word cloud for the top 50 words used in tweets sent from the iPhone. How do these word clouds compare? (Are there some common words frequently used from both sources? Are the most common words different between the sources?)

Don't forget to remove stop words before creating the word cloud. Also remove the terms "https" and "t.co".

ANSWER: The wordcloud generated by the Android tweets seem to be more focused on his opponents. The most common word being "Hillary," and other accompanying words such as "crooked" and "bad" imply that these tweets were created to negatively portray the other candidates. On the other hand, the iPhone tweets seemed to be more geared to progressing his campaign. While "Hillary" is also very common, common words in the iPhone tweets are "trump2016," "makeamericagreatagain," and "votetrump." Although there are many common words shared in both the Android and iPhone tweets, it appears that the main motivation behind the tweets differed.

```
data("stop_words")

# Add https and t.co to stop words list
custom_stop_words <- tribble(
  ~word, ~lexicon,
  "https", "CUSTOM",
  "t.co", "CUSTOM"
)

stop_words <- bind_rows(stop_words, custom_stop_words)

# Create wordcloud for Android
android <- tweets2 %>%
  filter(source == "Android") %>%
  rename(tweet = text) # can't replicate text column

android <- tibble(text = android$tweet) %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE)

android %>%
  head(50) %>%
  wordcloud2::wordcloud2(size = 0.5, color='random-dark')
```


2b. Create a visualization that compares the top 10 *bigrams* appearing in tweets by each source (that is, facet by source). After creating a dataset with one row per bigram, you should remove any rows that contain a stop word within the bigram.

How do the top used bigrams compare between the two sources?

ANSWER: Looking at the visualization, there are similar conclusions to be made as above. The top bigrams of the Android tweets are much more negative and targeted towards his opponents, such as “crooked hillary,” “lyin ted,” and “bad judgement.” On the other hand, the iPhone tweet bigrams are much more focused on gaining support for his campaign, with phrases such as “makeamericagreatagain trump2016,” “america safe,” “america trump2016,” “indiana trump2016,” “trump2016 americafirst,” and “trump2016 supertuesday.”

```
library(tidytext)

# Get bigrams of Trump's tweets
trump_bigrams <- tweets2 %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%
  group_by(source) %>%
  count(bigram, sort = TRUE)

# Remove stop words
bigrams_separated <- trump_bigrams %>%
  separate(bigram, c("word1", "word2"), sep = " ")

bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word)


trump_bigrams_final <- bigrams_filtered %>%
  unite(bigram, word1, word2, sep = " ") %>%
  filter(bigram != "https t.co") %>%
  arrange(desc(n)) %>%
  top_n(10)

head(trump_bigrams_final)

## # A tibble: 6 x 3
## # Groups:   source [2]
##   source bigram          n
##   <chr>   <chr>      <int>
## 1 Android crooked hillary    89
## 2 Android hillary clinton    54
## 3 iPhone  makeamericagreatagain trump2016 48
## 4 Android ted cruz          32
## 5 iPhone  crooked hillary    26
## 6 Android bernie sanders     22

# Visualization of top 10 bigrams across the different sources
trump_bigrams_final %>%
  ungroup %>%
  mutate(phrase = factor(bigram, levels = rev(unique(bigram)))) %>%
  ggplot(aes(phrase, n, fill = source)) +
```

```
geom_col(show.legend = FALSE) +
facet_wrap(~source, ncol = 2, scales = "free_y") +
labs(y = "Number of Bigrams", x = NULL) +
ggtitle("Top 10 Bigrams in Android and iPhone Trump Tweets") +
coord_flip()
```



2c. Consider the sentiment. Compute the proportion of words among the tweets within each source classified as “angry” and the proportion of words classified as “joy” based on the NRC lexicon. How does the proportion of “angry” and “joy” words compare between the two sources? What about “positive” and “negative” words?

ANSWER: Looking at the table below, there are not many instances of words that are classified as “anger” or “joy” in Trump’s tweets (for both iPhone and Android tweets both classifications consist of less than 6% of the total words). There are a few more words that are classified as anger than joy in the Android tweets, whereas the proportions are about the same in the iPhone tweets. In a more general sense, there is a larger proportion of positive and negative words in the Android tweets compared to the iPhone tweets.

```
nrc <- get_sentiments("nrc")

# Helper Function to Get Total Sentiment Tones from NRC Lexicon
total_sentiments <- function(nrc, word_count) {
  nrc %>%
    inner_join(word_count) %>%
    group_by(sentiment) %>%
    summarise(total = sum(n))
}

sentiments <- c("anger", "joy", "negative", "positive")

# Get sentiment proportions of Android
android_sentiments <- total_sentiments(nrc, android)
rownames(android_sentiments) <- android_sentiments$sentiment # to avoid using magic #s
android_sentiments <- c(android_sentiments[sentiments, ]$total/sum(android$n))

# Get sentiment proportions of iPhone
iphone_sentiments <- total_sentiments(nrc, iphone)
rownames(iphone_sentiments) <- iphone_sentiments$sentiment # to avoid using magic #s
iphone_sentiments <- c(iphone_sentiments[sentiments, ]$total/sum(iphone$n))

# Provide table of all the values
tibble(
  Sentiment = c("Anger", "Joy", "Positive", "Negative"),
  Android = c(android_sentiments),
  iPhone = c(iphone_sentiments)
)

## # A tibble: 4 x 3
##   Sentiment Android iPhone
##   <chr>      <dbl> <dbl>
## 1 Anger      0.0522 0.0369
## 2 Joy        0.0384 0.0354
## 3 Positive   0.0930 0.0569
## 4 Negative   0.105  0.0910
```

2d. Lastly, based on your responses above, do you think there is evidence to support Robinson's claim that Trump only writes the (angrier) Android half of the tweets from realDonaldTrump? In 2-4 sentences, please explain.

ANSWER: Robinson's claim is not entirely supported above. From parts a and b, it is clear that the majority of Android tweets are written to target his opponents, including words and phrases such as "crooked hillary," "lyin ted," "bad judgement," and "bad." However, we also see that the iPhone tweets constantly contain some of these phrases as well, which suggests that Trump writes angrier iPhone tweets as well. Furthermore, in part c, we see that the proportion of angry and negative words in Android tweets are very similar to those of the iPhone tweets. As such, although there is some evidence of Robinson's claim that the Android tweets are angrier, there also is evidence of angry iPhone tweets.